

OLAP Practice



KOREA UNIVERSITY
DATABASE LAB

hyubjinlee@korea.ac.kr

Contents

- OLAP Queries
 - GROUPING SETS
 - CUBE
 - ROLLUP
 - Window functions
- Practice Data
 - Data load
 - Table description
- Practice Exercise
- Homework

OLAP Queries

- GROUPING SETS
- CUBE
- ROLLUP
- Window functions

GROUPING SETS

- **GROUPING SETS** generates a result set equivalent to which generated by the **UNION ALL** of the multiple **GROUP BY** clauses

```
((GROUP BY expr1)
 UNION ALL
 (GROUP BY expr2)
 UNION ALL
 (GROUP BY expr3))
```



```
GROUPING SETS (expr1, expr2, expr3)
```

GROUPING SETS

```
SELECT customer_id, staff_id, sum(amount)  
FROM payment GROUP BY customer_id, staff_id;
```

UNION ALL

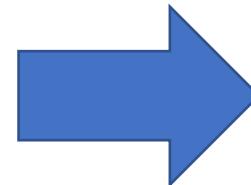
```
SELECT customer_id, NULL, sum(amount)  
FROM payment GROUP BY customer_id;
```

UNION ALL

```
SELECT NULL, staff_id, sum(amount)  
FROM payment GROUP BY staff_id;
```

UNION ALL

```
SELECT NULL, NULL, sum(amount)  
FROM payment;
```



```
SELECT customer_id, staff_id,  
sum(amount)  
FROM payment  
GROUP BY  
GROUPING SETS (  
    (customer_id, staff_id),  
    (customer_id),  
    (staff_id),  
    ()  
);
```

GROUPING SETS

```
SELECT customer_id, staff_id,  
       sum(amount)  
FROM payment  
GROUP BY  
    GROUPING SETS (  
        (customer_id, staff_id),  
        (customer_id),  
        (staff_id),  
        ()  
    );
```

customer_id smallint	staff_id smallint	sum numeric
448	2	76.83
459	1	108.78
460	1	46.90
236	2	94.80
282	2	52.87
112	1	56.07
⋮		⋮
449	[null]	00.00
64	[null]	91.70
520	[null]	127.69
55	[null]	84.81
148	[null]	211.55
[null]	1	30252.12
[null]	2	31059.92
[null]	[null]	61312.04

CUBE

- Analyze all possible subsets with more columns

GROUPING SETS (
(c1,c2,c3),
(c1,c2),
(c1,c3),
(c2,c3),
(c1),
(c2),
(c3),
(
)



CUBE(c1,c2,c3)

GROUPING SETS (
(c1,c2,c3,c4),
(c1,c2,c3),
...
(c1)
(c2),
(c3),
(c4),
(
)



CUBE(c1,c2,c3,c4)

CUBE

- PostgreSQL allows you to perform a partial cube to reduce the number of aggregates calculated.

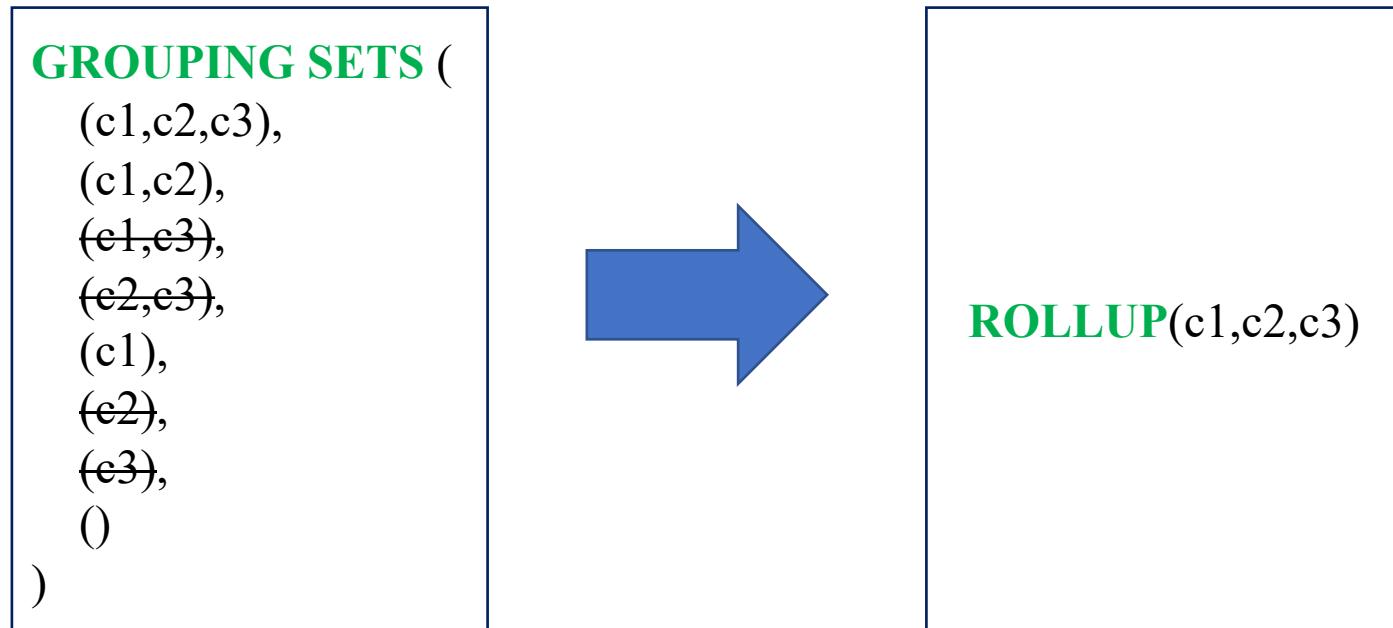
```
SELECT c1, c2, c3, aggregate (c4)  
FROM table_name  
GROUP BY c1, CUBE(c2, c3);
```



```
GROUPING SETS (  
    (c1,c2,c3),  
    (c1,c2),  
    (c1,c3),  
    (c1),  
)
```

ROLLUP

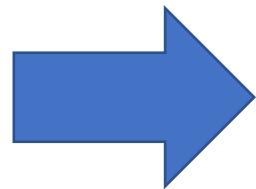
- ROLLUP(c1,c2,c3) generates only four grouping sets, **assuming the hierarchy c1 > c2 > c3** as follows



ROLLUP

- A common use of ROLLUP is to calculate the aggregations of data by year, month, and date, considering the hierarchy year > month > date
- You can extract (year, month, day, hour, minute and second) from *timestamp* data type through **EXTRACT**

```
dvdrental=# select rental_date from rental;  
rental_date  
-----  
2005-05-24 22:54:33  
2005-05-24 23:03:39  
2005-05-24 23:04:41  
2005-05-24 23:05:21  
2005-05-24 23:08:07  
2005-05-24 23:11:53  
2005-05-24 23:31:46  
2005-05-25 00:00:40  
2005-05-25 00:02:21  
2005-05-25 00:09:02  
2005-05-25 00:19:27
```



extract (year from rental_date)
Y
M
D
h
min
s

```
dvdrental=# SELECT  
dvdrental-# EXTRACT(year from rental_date) Y,  
dvdrental-# EXTRACT(month from rental_date) M,  
dvdrental-# EXTRACT(day from rental_date) D,  
dvdrental-# EXTRACT(hour from rental_date) h,  
dvdrental-# EXTRACT(minute from rental_date) min,  
dvdrental-# EXTRACT(second from rental_date) s  
dvdrental-# FROM rental;  
y | m | d | h | min | s  
---+---+---+---+---+---  
2005 | 5 | 24 | 22 | 54 | 33  
2005 | 5 | 24 | 23 | 3 | 39  
2005 | 5 | 24 | 23 | 4 | 41  
2005 | 5 | 24 | 23 | 5 | 21  
2005 | 5 | 24 | 23 | 8 | 7
```

Window Functions

[Input 데이터가 다른 행을 사용하는!
그것은 윈도우]

- A window function performs a calculation across a set of table rows that are somehow related to the current row.
- This is comparable to the type of calculation that can be done with an aggregate function.
- But unlike regular aggregate functions, use of a window function does not cause rows to become grouped into a single output row — the rows retain their separate identities.

optional [**window function** (arg1, arg2,..) OVER (*누적되는 행들, 결과*
[**PARTITION BY** partition_expression]
[**ORDER BY** sort_expression [ASC | DESC] [NULLS {FIRST | LAST }]
)
(partition은 order)]

- The **PARTITION BY** list within **OVER** specifies dividing the row into groups, or partitions, that share the same values of the **PARTITION BY** expression(s)

An Example Relation

Instructor Table

id [PK] character varying (5)	name character varying (20)	dept_name character varying (20)	salary numeric (8,2)
10101	Srinivasan	Comp. Sci.	65000.00
12121	Wu	Finance	90000.00
15151	Mozart	Music	40000.00
22222	Einstein	Physics	95000.00
32343	El Said	History	60000.00
33456	Gold	Physics	87000.00
45565	Katz	Comp. Sci.	75000.00
58583	Califieri	History	62000.00
76543	Singh	Finance	80000.00
76766	Crick	Biology	72000.00
83821	Brandt	Comp. Sci.	92000.00
98345	Kim	Elec. Eng.	80000.00
32134	Silver	Comp. Sci.	75000.00

PARTITION BY

- For each row, the window function is computed across the rows that fall into the same partition as the current row.

Select dept_name, id, name, salary, sum(salary) over (partition by dept_name)

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name	id	name	salary	sum
	[PK]		numeric (8,2)	numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	32134	Silver	75000.00	307000.00
Comp. Sci.	45565	Katz	75000.00	307000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Comp. Sci.	10101	Srinivasan	65000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	170000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	122000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	22222	Einstein	95000.00	182000.00
Physics	33456	Gold	87000.00	182000.00



PARTITION BY

- The **PARTITION BY** clause is optional. If you skip the **PARTITION BY** clause, the window function will treat the whole result set as a single partition.

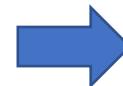
```
SELECT dept_name, id, name, salary,  
sum(salary) OVER ()  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Comp. Sci.	10101	Srinivasan	65000.00	973000.00
Finance	12121	Wu	90000.00	973000.00
Music	15151	Mozart	40000.00	973000.00
Physics	22222	Einstein	95000.00	973000.00
History	32343	El Said	60000.00	973000.00
Physics	33456	Gold	87000.00	973000.00
Comp. Sci.	45565	Katz	75000.00	973000.00
History	58583	Califieri	62000.00	973000.00
Finance	76543	Singh	80000.00	973000.00
Biology	76766	Crick	72000.00	973000.00
Comp. Sci.	83821	Brandt	92000.00	973000.00
Elec. Eng.	98345	Kim	80000.00	973000.00
Comp. Sci.	32134	Silver	75000.00	973000.00

Aggregation Functions vs. Window Functions

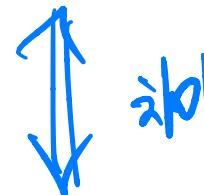
Aggregation function

```
SELECT dept_name, sum(salary)  
FROM instructor  
GROUP BY dept_name  
ORDER BY dept_name;
```



dept_name	sum
Biology	72000.00
Comp. Sci.	307000.00
Elec. Eng.	80000.00
Finance	170000.00
History	122000.00
Music	40000.00
Physics	182000.00

⇒ 'dept-' 7개가 있음



Window function

```
SELECT dept_name,  
sum(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

over (Partition By
dept_name)



dept_name	sum
Biology	72000.00
Comp. Sci.	307000.00
Elec. Eng.	80000.00
Finance	170000.00
Finance	170000.00
History	122000.00
History	122000.00
Music	40000.00
Physics	182000.00
Physics	182000.00

⇒ 'custom-' Record 7개가 있음

Why Window Functions ?

- Window function is useful when you want to know the individual value and entire group value (sum, avg, ...) at the same time

```
SELECT dept_name, id, name, salary,
       sum(salary) OVER (PARTITION BY dept_name)
  FROM instructor;
```

dept_name				salary	sum
character varying (20)	[PK]	character varying (5)		numeric (8,2)	numeric
Biology	76766	Crick		72000.00	72000.00
Comp. Sci.	32134	Silver		75000.00	307000.00
Comp. Sci.	45565	Katz		75000.00	307000.00
Comp. Sci.	83821	Brandt		92000.00	307000.00
Comp. Sci.	10101	Srinivasan		65000.00	307000.00
Elec. Eng.	98345	Kim		80000.00	80000.00
Finance	76543	Singh		80000.00	170000.00
Finance	12121	Wu		90000.00	170000.00
History	32343	El Said		60000.00	122000.00
History	58583	Califieri		62000.00	122000.00
Music	15151	Mozart		40000.00	40000.00
Physics	22222	Einstein		95000.00	182000.00
Physics	33456	Gold		87000.00	182000.00

Multiple Window Functions

- You can use multiple window functions

```
SELECT  
    wf1() OVER(PARTITION BY c1 ORDER BY c2),  
    wf2() OVER(PARTITION BY c3 ORDER BY c4)  
FROM table_name;
```

Multiple Window Functions

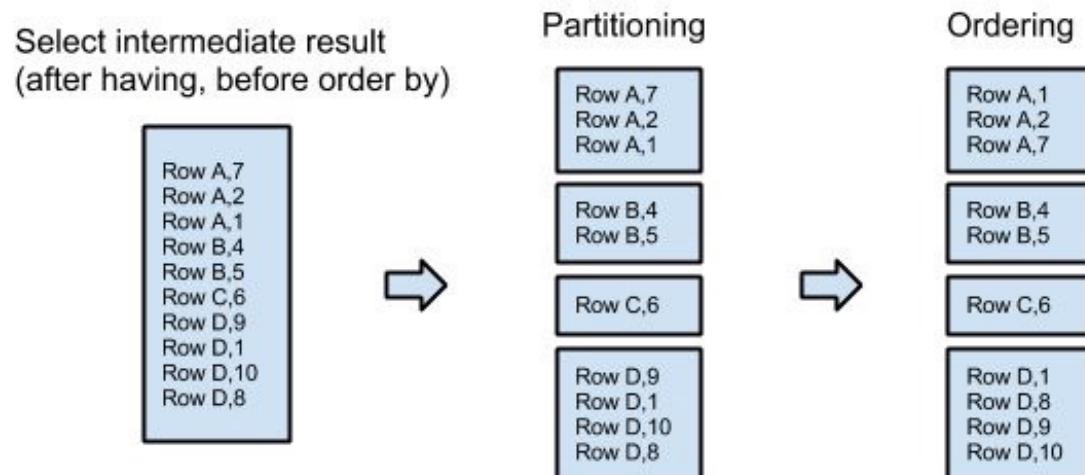
```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name),  
avg(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name character varying (20)	id [PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric	avg numeric
Biology	76766	Crick	72000.00	72000.00	72000.00000000000000
Comp. Sci.	32134	Silver	75000.00	307000.00	76750.00000000000000
Comp. Sci.	45565	Katz	75000.00	307000.00	76750.00000000000000
Comp. Sci.	83821	Brandt	92000.00	307000.00	76750.00000000000000
Comp. Sci.	10101	Srinivasan	65000.00	307000.00	76750.00000000000000
Elec. Eng.	98345	Kim	80000.00	80000.00	80000.00000000000000
Finance	76543	Singh	80000.00	170000.00	85000.00000000000000
Finance	12121	Wu	90000.00	170000.00	85000.00000000000000
History	32343	El Said	60000.00	122000.00	61000.00000000000000
History	58583	Califieri	62000.00	122000.00	61000.00000000000000
Music	15151	Mozart	40000.00	40000.00	40000.00000000000000
Physics	22222	Einstein	95000.00	182000.00	91000.00000000000000
Physics	33456	Gold	87000.00	182000.00	91000.00000000000000

ORDER BY

- You can also control the order in which rows are processed by window functions using **ORDER BY** within **OVER**

```
window_function(arg1, arg2,..) OVER (  
    [PARTITION BY partition_expression] (이 줄은 안쓰면 안되네!)  
    [ORDER BY sort_expression [ASC | DESC] [NULLS {FIRST | LAST} ]]  
)
```



ORDER BY

- You can also control the order in which rows are processed by window functions using **ORDER BY** within **OVER**

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	215000.00
Comp. Sci.	45565	Katz	75000.00	215000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	182000.00

Window Frame

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

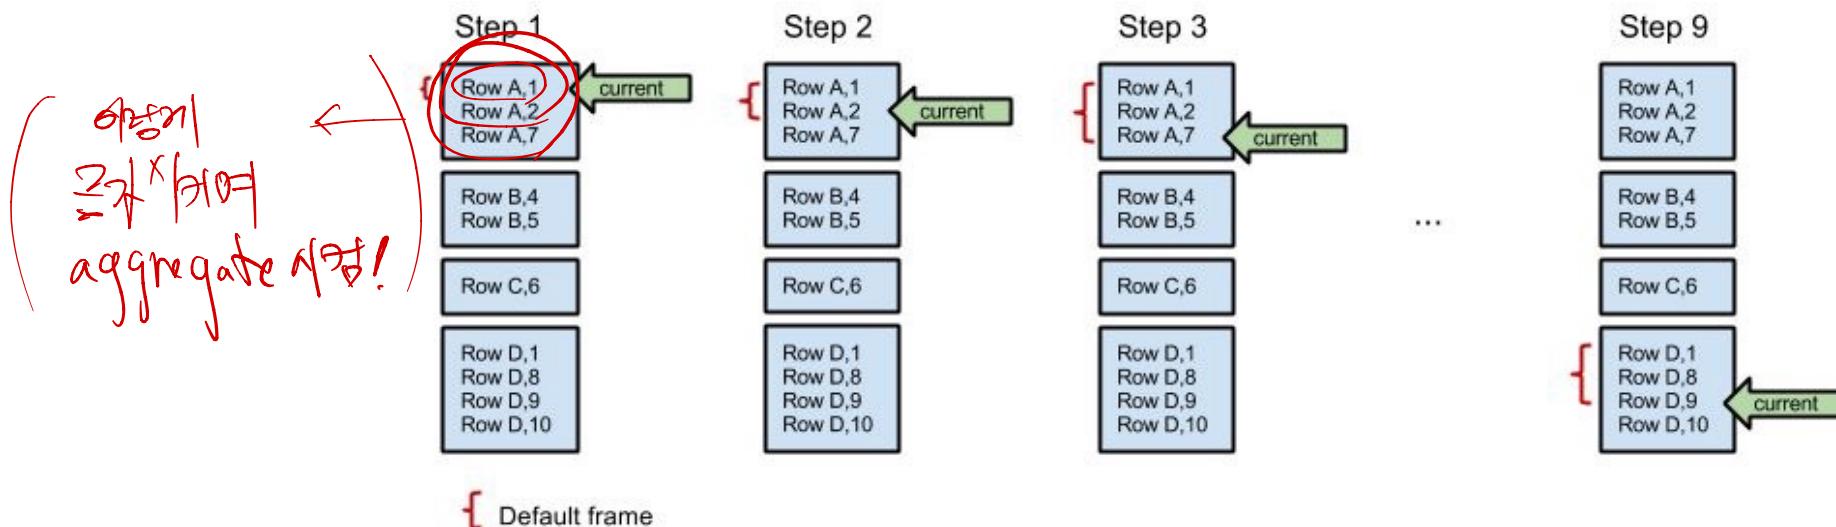
dept_name character varying (20)	[PK] id character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	215000.00
Comp. Sci.	45565	Katz	75000.00	215000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	182000.00

Why sum is different in
the same group?

Partition by
order by 2 3 1 4

Window Frame

- There is another important concept associated with window functions: for each row, there is a set of rows within its partition called its window frame
 - Many (but not all) window functions act only on the rows of the window frame, rather than of the whole partition
 - $\text{avg}()$, $\text{min}()$, $\text{max}()$, $\text{sum}()$, $\text{count}()$ etc.
- When **ORDER BY** is omitted the default frame consists of all rows in the partition
- By default, if **ORDER BY** is supplied then the frame consists of all rows from the start of the partition up through the current row, plus any following rows that are equal to the current row according to the **ORDER BY** clause



Window Frame without ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	32134	Silver	75000.00	307000.00
Comp. Sci.	45565	Katz	75000.00	307000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Comp. Sci.	10101	Srinivasan	65000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	170000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	122000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	22222	Einstein	95000.00	182000.00
Physics	33456	Gold	87000.00	182000.00

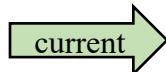
current
Each row
Total sum
Sum of all rows

Window frame

$$\text{Sum} = 65000 + 2 \cdot 75000 + 92000 \\ = 307000$$

Window Frame without ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```



dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	32134	Silver	75000.00	307000.00
Comp. Sci.	45565	Katz	75000.00	307000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Comp. Sci.	10101	Srinivasan	65000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	170000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	122000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	22222	Einstein	95000.00	182000.00
Physics	33456	Gold	87000.00	182000.00

Window frame Sum = $65000 + 2 \cdot 75000 + 92000 = 307000$

Window Frame without ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	32134	Silver	75000.00	307000.00
Comp. Sci.	45565	Katz	75000.00	307000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Comp. Sci.	10101	Srinivasan	65000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	170000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	122000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	22222	Einstein	95000.00	182000.00
Physics	33456	Gold	87000.00	182000.00

current →

Window frame

Sum = $65000 + 2*75000 + 92000$
 $= 307000$

Window Frame without ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	32134	Silver	75000.00	307000.00
Comp. Sci.	45565	Katz	75000.00	307000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Comp. Sci.	10101	Srinivasan	65000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	170000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	122000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	22222	Einstein	95000.00	182000.00
Physics	33456	Gold	87000.00	182000.00

Window frame

$$\text{Sum} = 65000 + 2*75000 + 92000 \\ = 307000$$

Window Frame with ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	215000.00
Comp. Sci.	45565	Katz	75000.00	215000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	182000.00

Window Frame with ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	215000.00
Comp. Sci.	45565	Katz	75000.00	215000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	182000.00

current →

Window frame

$$\text{Sum} = 65000 + 2 \cdot 75000 \\ = 215000$$

Window Frame with ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	215000.00
Comp. Sci.	45565	Katz	75000.00	215000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	182000.00

current →

Window frame

$$\text{Sum} = 65000 + 2 \cdot 75000 \\ = 215000$$

Window Frame with ORDER BY

```
SELECT dept_name, id, name, salary,  
sum(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	sum numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	215000.00
Comp. Sci.	45565	Katz	75000.00	215000.00
Comp. Sci.	83821	Brandt	92000.00	307000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	170000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	122000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	182000.00

current →

Window frame

Sum = $65000 + 2*75000 + 92000$
 $= 307000$

Window Frame

- You can also create your own window frame
- But out of our range...

The diagram illustrates the execution of a window frame query across three stages. Each stage shows a table with data rows numbered 1 to 7. A 'current row' is indicated by a light gray background. In the first stage, row 4 is the current row. In the second stage, row 4 has moved to the preceding row position, and row 5 is now the current row. In the third stage, row 5 has moved to the preceding row position, and row 6 is now the current row. Arrows labeled 'preceding row' point to the previous row in each stage, and arrows labeled 'following row' point to the next row.

1	2015	7.00
2	2015	7.50
3	2015	8.00
4	2015	8.30
5	2015	8.50
6	2015	8.80
7	2015	9.00

current row →

1	2015	7.00
2	2015	7.50
3	2015	8.00
4	2015	8.30
5	2015	8.50
6	2015	8.80
7	2015	9.00

current row →

1	2015	7.00
2	2015	7.50
3	2015	8.00
4	2015	8.30
5	2015	8.50
6	2015	8.80
7	2015	9.00

current row →

← preceding row

← following row

Window Functions

- AVG(), MIN(), MAX(), SUM() and COUNT()
- ROW_NUMBER()
 - Number the current row within its partition starting from 1
- FIRST_VALUE()
 - Return a value evaluated against the first row within its partition
- LAST_VALUE()
 - Return a value evaluated against the last row within its partition
- NTH_VALUE()
 - Return a value evaluated against the nth row in an ordered partition
- RANK()
 - Rank the current row within its partition with gaps
- DENSE_RANK()
 - Rank the current row within its partition without gaps

7 1 2 3 4
↓ ↗ ↗ ↗ ↗
1 2 2 2 3 --

Window Functions

Sum, avg, count etc! row_number()
rank()
dense_rank()

- The **ROW_NUMBER()** function assigns a sequential number to each row in each partition.

```
SELECT dept_name, id, name, salary,  
ROW_NUMBER() OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	row_number bigint
Biology	76766	Crick	72000.00	1
Comp. Sci.	10101	Srinivasan	65000.00	1
Comp. Sci.	32134	Silver	75000.00	2
Comp. Sci.	45565	Katz	75000.00	3
Comp. Sci.	83821	Brandt	92000.00	4
Elec. Eng.	98345	Kim	80000.00	1
Finance	76543	Singh	80000.00	1
Finance	12121	Wu	90000.00	2
History	32343	El Said	60000.00	1
History	58583	Califieri	62000.00	2
Music	15151	Mozart	40000.00	1
Physics	33456	Gold	87000.00	1
Physics	22222	Einstein	95000.00	2

Window Functions

```
SELECT dept_name, id, name, salary,  
ROW_NUMBER() OVER (PARTITION BY dept_name)  
FROM instructor;
```



dept_name character varying (20)	[PK] id character varying (5)	name character varying (20)	salary numeric (8,2)	row_number bigint
Biology	76766	Crick	72000.00	1
Comp. Sci.	32134	Silver	75000.00	1
Comp. Sci.	45565	Katz	75000.00	2
Comp. Sci.	83821	Brandt	92000.00	3
Comp. Sci.	10101	Srinivasan	65000.00	4
Elec. Eng.	98345	Kim	80000.00	1
Finance	76543	Singh	80000.00	1
Finance	12121	Wu	90000.00	2
History	32343	El Said	60000.00	1
History	58583	Califieri	62000.00	2
Music	15151	Mozart	40000.00	1
Physics	22222	Einstein	95000.00	1
Physics	33456	Gold	87000.00	2

Window Functions

first_value(salary)

- **FIRST_VALUE()** function returns a value evaluated against the first row within its partition, whereas the **LAST_VALUE()** function returns a value evaluated against the last row in its partition.
- The following statement uses the **FIRST_VALUE()** to return the lowest price for every product group.

```
SELECT dept_name, id, name, salary,  
FIRST_VALUE(salary) OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	first_value numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	10101	Srinivasan	65000.00	65000.00
Comp. Sci.	32134	Silver	75000.00	65000.00
Comp. Sci.	45565	Katz	75000.00	65000.00
Comp. Sci.	83821	Brandt	92000.00	65000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	80000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	60000.00
Music	15151	Mozart	40000.00	40000.00
Physics	33456	Gold	87000.00	87000.00
Physics	22222	Einstein	95000.00	87000.00

first_value

~~sum(salary)~~

over(partition by
_____?)

order by _____

Window Functions

```
SELECT dept_name, id, name, salary,  
FIRST_VALUE(salary) OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name character varying (20)	id [PK] character varying (5)	name character varying (20)	salary numeric (8,2)	first_value numeric
Biology	76766	Crick	72000.00	72000.00
Comp. Sci.	32134	Silver	75000.00	75000.00
Comp. Sci.	45565	Katz	75000.00	75000.00
Comp. Sci.	83821	Brandt	92000.00	75000.00
Comp. Sci.	10101	Srinivasan	65000.00	75000.00
Elec. Eng.	98345	Kim	80000.00	80000.00
Finance	76543	Singh	80000.00	80000.00
Finance	12121	Wu	90000.00	80000.00
History	32343	El Said	60000.00	60000.00
History	58583	Califieri	62000.00	60000.00
Music	15151	Mozart	40000.00	40000.00
Physics	22222	Einstein	95000.00	95000.00
Physics	33456	Gold	87000.00	95000.00

Window Functions

- **RANK()** function assigns a rank to each row within an ordered partition.

```
SELECT dept_name, id, name, salary,  
RANK() OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	id [PK] character varying (5)	name character varying (20)	salary numeric (8,2)	rank bigint
Biology	76766	Crick	72000.00	1
Comp. Sci.	10101	Srinivasan	65000.00	1
Comp. Sci.	32134	Silver	75000.00	2
Comp. Sci.	45565	Katz	75000.00	2
Comp. Sci.	83821	Brandt	92000.00	4
Elec. Eng.	98345	Kim	80000.00	1
Finance	76543	Singh	80000.00	1
Finance	12121	Wu	90000.00	2
History	32343	El Said	60000.00	1
History	58583	Califieri	62000.00	2
Music	15151	Mozart	40000.00	1
Physics	33456	Gold	87000.00	1
Physics	22222	Einstein	95000.00	2

Window Functions

```
SELECT dept_name, id, name, salary,  
RANK() OVER (PARTITION BY dept_name)  
FROM instructor;
```

dept_name character varying (20)	[PK] id character varying (5)	name character varying (20)	salary numeric (8,2)	rank bigint
Biology	76766	Crick	72000.00	1
Comp. Sci.	32134	Silver	75000.00	1
Comp. Sci.	45565	Katz	75000.00	1
Comp. Sci.	83821	Brandt	92000.00	1
Comp. Sci.	10101	Srinivasan	65000.00	1
Elec. Eng.	98345	Kim	80000.00	1
Finance	76543	Singh	80000.00	1
Finance	12121	Wu	90000.00	1
History	32343	El Said	60000.00	1
History	58583	Califieri	62000.00	1
Music	15151	Mozart	40000.00	1
Physics	22222	Einstein	95000.00	1
Physics	33456	Gold	87000.00	1

order by id
rank 1/5/1/2
partition
at 50% 1/2
dept - ordered
(sp15ch22ptor2)

Window Functions

- **DENSE_RANK()** function assigns a rank to each row within an ordered partition, but the ranks have no gap. In other words, the same ranks are assigned to multiple rows and no ranks are skipped.

```
SELECT dept_name, id, name, salary,  
DENSE_RANK() OVER (PARTITION BY dept_name ORDER BY salary)  
FROM instructor;
```

dept_name character varying (20)	[PK] character varying (5)	name character varying (20)	salary numeric (8,2)	dense_rank bigint
Biology	76766	Crick	72000.00	1
Comp. Sci.	10101	Srinivasan	65000.00	1
Comp. Sci.	32134	Silver	75000.00	2
Comp. Sci.	45565	Katz	75000.00	2
Comp. Sci.	83821	Brandt	92000.00	3
Elec. Eng.	98345	Kim	80000.00	1
Finance	76543	Singh	80000.00	1
Finance	12121	Wu	90000.00	2
History	32343	El Said	60000.00	1
History	58583	Califieri	62000.00	2
Music	15151	Mozart	40000.00	1
Physics	33456	Gold	87000.00	1
Physics	22222	Einstein	95000.00	2

Other Window Functions

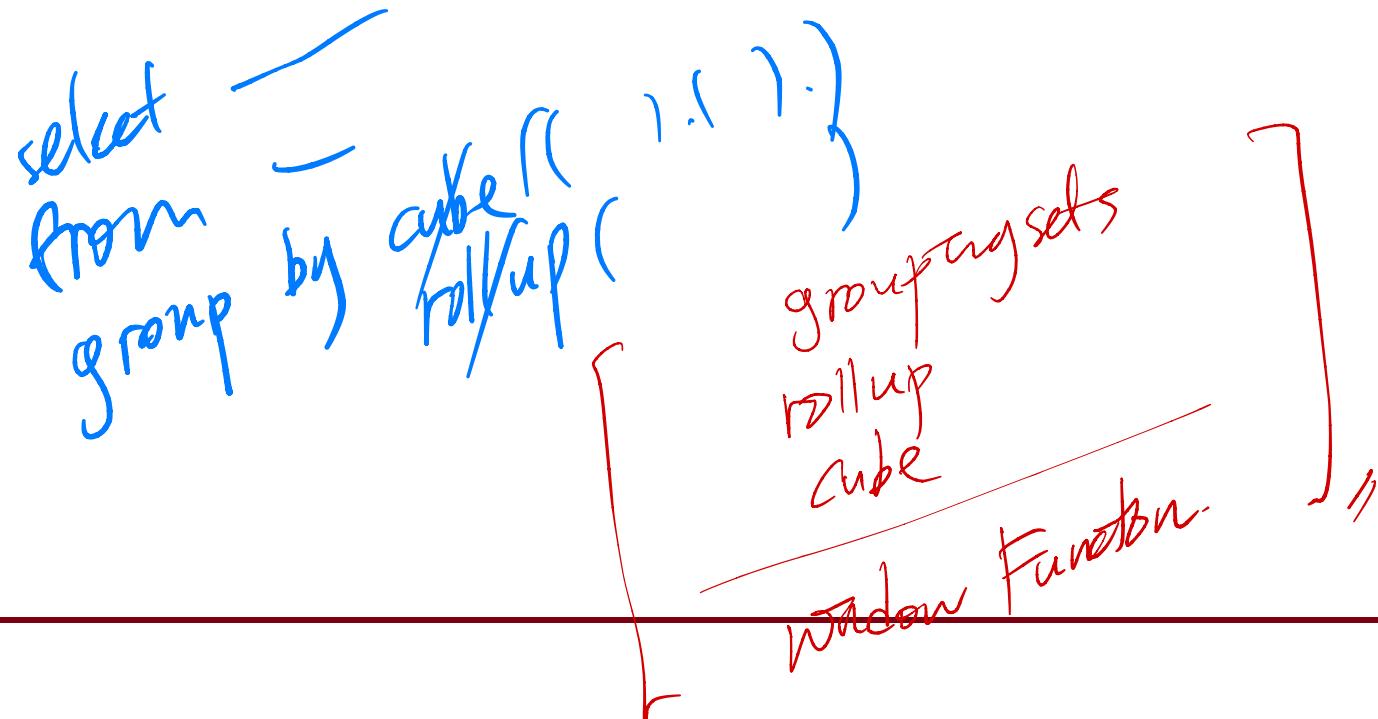
- <http://www.postgresqltutorial.com/postgresql-window-function/>

Name	Description
CUME_DIST	Return the relative rank of the current row.
DENSE_RANK	Rank the current row within its partition without gaps.
FIRST_VALUE	Return a value evaluated against the first row within its partition.
LAG	Return a value evaluated at the row that is at a specified physical offset row before the current row within the partition.
LAST_VALUE	Return a value evaluated against the last row within its partition.
LEAD	Return a value evaluated at the row that is offset rows after the current row within the partition.
NTILE	Divide rows in a partition as equally as possible and assign each row an integer starting from 1 to the argument value.
NTH_VALUE	Return a value evaluated against the nth row in an ordered partition.

Practice Data

- Data load
- Table description

select
from
group by) cube
rollup () . .
grouping sets
rollup
cube
Wacion Function



Data Load

1. Download files from blackboard
 - dvdrental.tar
 - remove_lastupdate.sql
2. Run *psql*
3. Create a new database '*CREATE DATABASE dvdrental;*'

Data Load

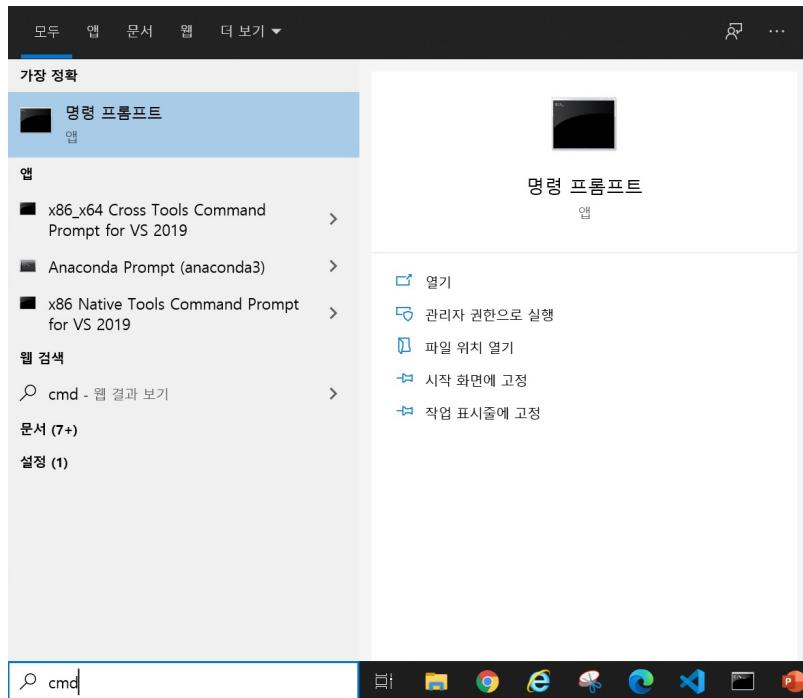
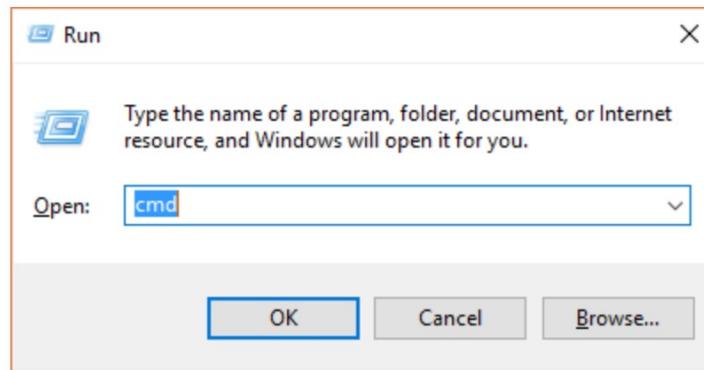
- Run psql and check the psql version

```
SQL Shell (psql)
Server [localhost]:
Database [postgres]:
Port [5432]:
Username [postgres]:
postgres 사용자의 암호:
psql (13.2)
도움말을 보려면 "help"를 입력하십시오.

postgres=#
```

Data Load

- Run command prompt
 - Window key + R
 - Search cmd on window



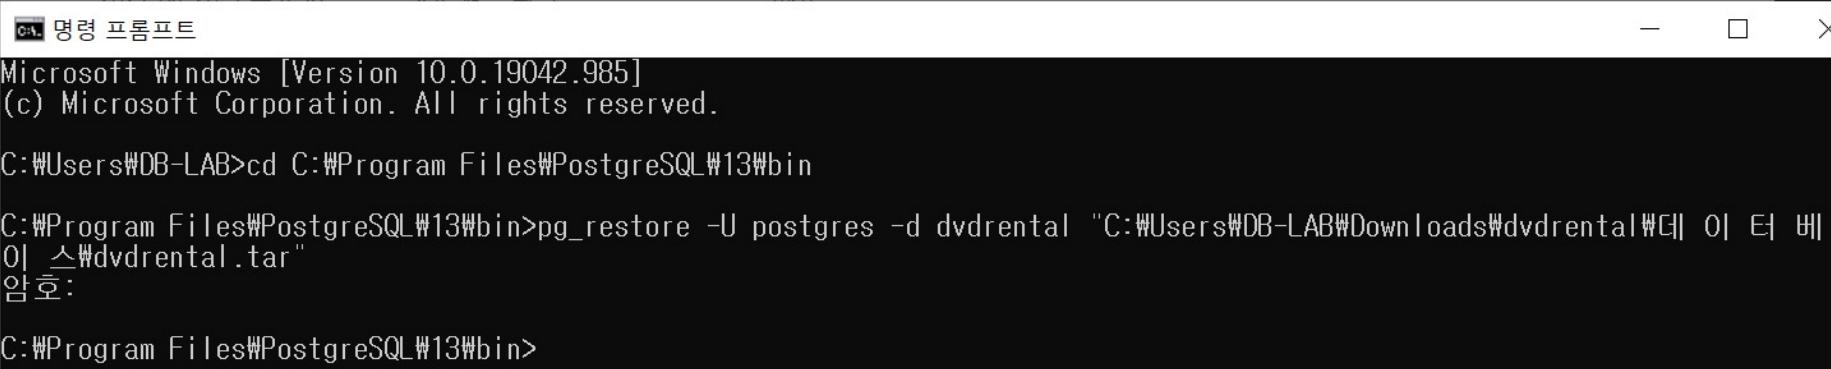
Data Load

- `cd C:\Program Files\PostgreSQL\13\bin`
 - It is default path of postgresql 13
 - If it does not work, check your installation path of postgresql or version
- `pg_restore -U postgres -d dvdrental "DVDRENTAL.TAR PATH"`
 - If whitespace is included in the path, use **double** quotation("")
 - When entering password, check the CapsLock, and language(한/영)

주의사항:

'cmd' 창에서 실행함.

(postgresql에서 실행하는거 아님)



```
명령 프롬프트
Microsoft Windows [Version 10.0.19042.985]
(c) Microsoft Corporation. All rights reserved.

C:\Users\DB-LAB>cd C:\Program Files\PostgreSQL\13\bin

C:\Program Files\PostgreSQL\13\bin>pg_restore -U postgres -d dvdrental "C:\Users\DB-LAB\Downloads\dvdrental\데이터베이스\스냅샷\dvdrental.tar"
암호:

C:\Program Files\PostgreSQL\13\bin>
```

Data Load

- connect dydrental database

- \c dvdrental

- execute remove_lastupdate.sql script file

- \i 'sql script file path'

- If whitespace or Korean is included in the path, use **single** quotation(')
 - Be sure that file separator is double-backslash(\) or slash(/)

```
alter table category drop column last_update;
alter table inventory drop column last_update;
alter table customer drop column last_update;
alter table film_category drop column last_update;
alter table rental drop column last_update;
alter table address drop column last_update;
alter table film drop column last_update;
alter table staff drop column last_update;
alter table city drop column last_update;
alter table language drop column last_update;
alter table country drop column last_update;
alter table film_actor drop column last_update;
alter table actor drop column last_update;
alter table store drop column last_update;
```

Table Description

Crow Foot Notation Symbols

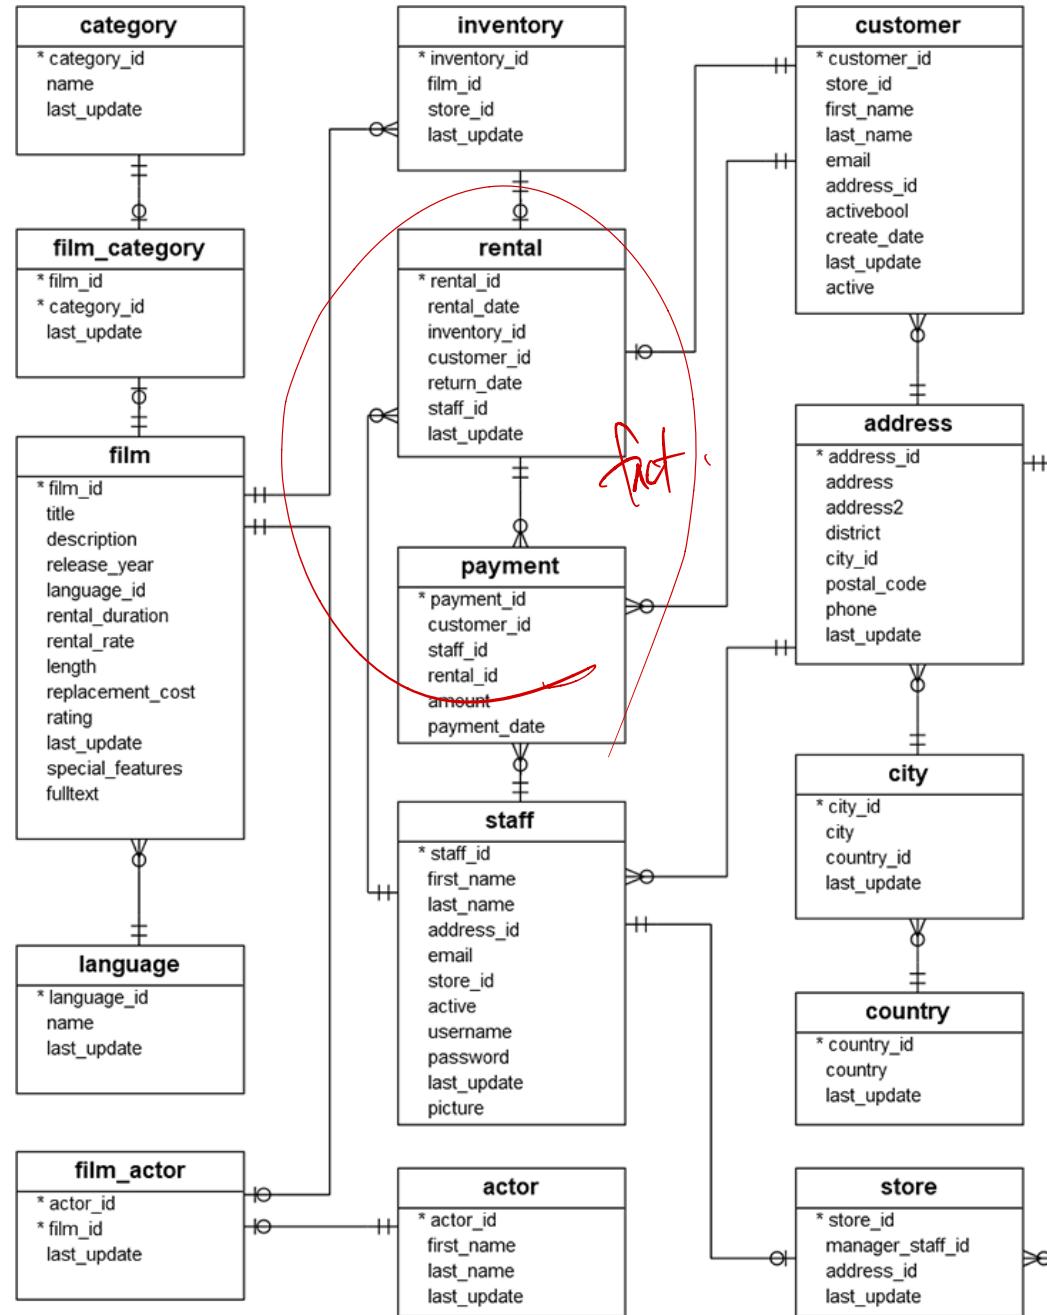
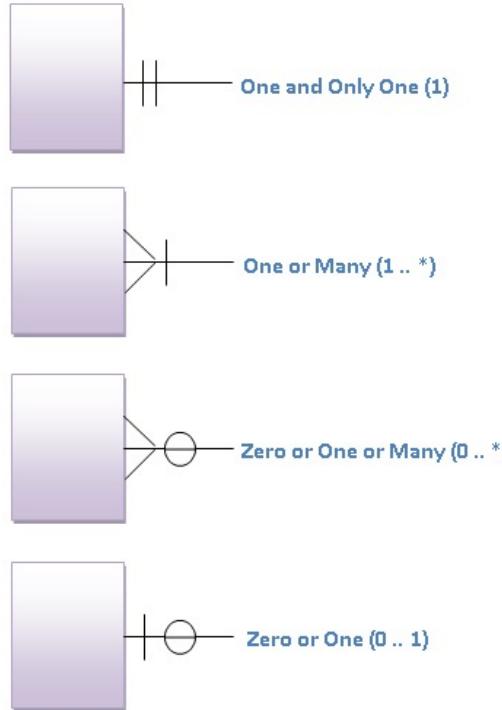


Table Description

- actor — contains actors data including first name and last name.
- film — contains films data such as title, release year, length, rating, etc.
- film_actor — contains the relationships between films and actors.
- category — contains film's categories data.
- film_category — containing the relationships between films and categories.
- store — contains the store data including manager staff and address.
- inventory — stores inventory data.
- rental — stores rental data.
- payment — stores customer's payments.
- staff — stores staff data.
- customer — stores customer's data.
- address — stores address data for staff and customers
- city — stores the city names.
- country — stores the country names.

Practice Exercise

- English version
- Korean version

Practice Exercise (English version)

Answer the following questions in SQL

1. Use *grouping sets* to count the number of films for each *rental_rate* and *rating*, respectively
2. Print the information about all combinations of *actor_id* and *category* name on the number of films each actor with *actor_id* 1 and 2 starred in (doesn't mean that they starred at the same time)
3. Use *rollup* to show total amount on *rental_date* by *year*, *month*, and *day*
4. For each film *category*, find *customer_ids* who is TOP-2 in order of the number of DVDs rented (if there are many ties, more than 2 customers can be printed)
5. Print the total amount for each customer, the total amount for each country, and *dense ranking* by country for total amount (Hint: First calculate the customer's total amount using the *with* clause, and then apply a *window query*)

실습 문제 (한국어)

다음을 SQL 문장으로 작성하시오

1. *grouping sets*를 사용하여 rental_rate 및 rating 각각에 대한 영화의 수를 출력하세요
2. actor_id가 1과 2인 두 배우들이 출연한 영화에 대해 (두 배우가 동시에 출연한 영화를 의미하는 것은 아님), actor_id와 category name의 모든 조합별로 각 배우가 출연한 영화의 수를 출력하세요
3. *rollup*을 사용하여 rental_date의 연도별, 월별, 일별 총 amount를 출력하세요
4. 영화의 각 카테고리별로, 대여한 DVD의 수가 많은 순으로 2위까지의 customer_id들을 찾습니다. (동률이 많으면 2명보다 많은 고객이 출력될 수 있음)
5. 고객별 총 amount, 나라별 총 amount, 나라별 총 amount에 대한 *dense ranking*을 출력하세요 (힌트: *with* 절을 사용하여 먼저 고객별 총 amount를 계산한 다음에 *window function*을 사용함)

Homework

1. Complete today's practice exercise
2. Take screenshots of your queries and execution results
3. Submit your report on blackboard
 - Before the next class hour
 - **Only PDF files** are accepted
 - **No late submission**